

CMPT 770/479: Distributed Systems (Fall 2021)

Assignment 3 - Report

Instructions:

- This report is worth 45 points.
- Answer in the space provided. Answers spanning beyond 3 lines (11pt font) will lose points.
- Input graphs used are available at the following location.
 - live-journal graph (LJ graph): `/scratch/input_graphs/lj`
 - RMAT graph: `/scratch/input_graphs/rmat`
- All your answers must be based on the experiments conducted with 4 workers on slow nodes. Answers based on fast nodes and/or different numbers of workers will result in 0 points.
- All the times are in seconds.

-
1. [12 points] Run Triangle Counting with `--strategy=1` on the LJ graph and the RMAT graph. Update the thread statistics in the tables below. What is your observation on the difference in time taken by each thread for RMAT and that for LJ? Why does this happen?

Answer:

The LJ graph is much more dense than the RMAT graph, this can be seen by the (almost equal) number of edges and vertices across the threads in strategy 1, so the workload is almost identical. Whereas, in LJ graph the workload is not divided equally so there is a divide in the thread timings.

Triangle Counting on LJ: Total time = 53.81312 seconds.

| thread_id | num_vertices | num_edges | triangle_count | time_taken |
|-----------|--------------|-----------|----------------|------------|
| 0 | 1211892 | 42920131 | 339204160 | 53.812647 |
| 1 | 1211892 | 15515692 | 213398914 | 11.230427 |
| 2 | 1211892 | 7141449 | 84872361 | 3.457746 |
| 3 | 1211895 | 3416501 | 45558451 | 0.979338 |

Triangle Counting on RMAT: Total time = 3.18986 seconds.

| thread_id | num_vertices | num_edges | triangle_count | time_taken |
|-----------|--------------|-----------|----------------|------------|
| 0 | 6249999 | 12650749 | 7 | 3.189352 |
| 1 | 6249999 | 12546666 | 5 | 3.156983 |
| 2 | 6249999 | 12399926 | 6 | 3.128009 |
| 3 | 6250002 | 12402659 | 9 | 3.121901 |

2. [9 points] Run Triangle Counting with `--strategy=2` on LJ graph. Update the thread statistics in the table below. Partitioning time is the time spent on task decomposition as required by `--strategy=2`. What is your observation on the difference in time taken by each thread, and the difference in `num_edges` for each thread? Are they correlated (yes/no)? Why?

Answer:

This time around the number of edges are much more comparable, and it can be seen that the time taken by each thread is much less than strategy 1 but the divide is still present in the timings of the threads, it might be because the actual counting is still done only on the basis of number of vertices.

Triangle Counting on LJ: Partitioning time = 0.0 seconds. Total time = 28.01175 seconds.

| thread_id | num_vertices | num_edges | triangle_count | time_taken |
|-----------|--------------|-----------|----------------|------------|
| 0 | 325540 | 17248443 | 136034326 | 27.966894 |
| 1 | 510545 | 17248443 | 144912885 | 21.973001 |
| 2 | 904041 | 17248443 | 219730109 | 15.745213 |
| 3 | 3107444 | 17248444 | 182356566 | 8.891030 |

3. [9 points] Run PageRank with `--strategy=1` on LJ graph. Update the thread statistics in the table below. What is your observation on the difference in time taken by each thread, and the difference in `num_edges` for each thread? Is the work uniformly distributed across threads (yes/no)? Why?

Answer:

Here it can be seen that the divide between the number of edges is a lot, but the time taken by the threads are comparable, that is because of the barrier waiting. The work is not uniformly divided across the threads, as the number of edges have a difference of order of 10 across the threads.

PageRank on LJ: Total time = 41.94746 seconds.

| thread_id | num_vertices | num_edges | time_taken |
|-----------|--------------|-----------|------------|
| 0 | 24237840 | 858402620 | 41.937171 |
| 1 | 24237840 | 310313840 | 41.937136 |
| 2 | 24237840 | 142828980 | 41.937097 |
| 3 | 24237900 | 68330020 | 41.922869 |

4. [9 points] Run PageRank with `--strategy=1` on LJ graph. Obtain the cumulative time spent by each thread on `barrier1` and `barrier2` (refer pagerank pseudocode for program 3 on assignment webpage) and update the table below. What is your observation on the difference in `barrier1_time` for each thread and the difference in `num_edges` for each thread? Are they correlated (yes/no)? Why?

Answer:

It can be seen that the 0th thread has very little wait time on the barriers because it has the most work to do, the rest of the threads wait on 0th thread before exiting the thread function. First for loop deals with the edges so barrier 1 time is higher for the threads > 0 (0th thread has most edges to compute).

PageRank on LJ: Total time = 41.94746 seconds.

| thread_id | num_vertices | num_edges | barrier1_time | barrier2_time | time_taken |
|-----------|--------------|-----------|---------------|---------------|------------|
| 0 | 24237840 | 858402620 | 0.000727 | 0.014058 | 41.937171 |
| 1 | 24237840 | 310313840 | 21.909496 | 0.010574 | 41.937136 |
| 2 | 24237840 | 142828980 | 31.171523 | 0.007112 | 41.937097 |
| 3 | 24237900 | 68330020 | 36.604755 | 0.004408 | 41.922869 |

5. [6 points] Run PageRank with `--strategy=2` on the LJ graph. Update the thread statistics in the table below. Update the time taken for task decomposition as required by `--strategy=2`. What is your observation on `barrier2_time` compared to the `barrier2_time` in question 4 above? Why are they same/different?

Answer:

Barrier 2 time is calculated after the for loop which runs on the vertices, and in edge distribution strategy there is a divide between the vertices and it is of 10 order across the threads, so the last thread has the most vertices, so its barrier 2 time is low so the other threads have to wait for barrier 2.

PageRank on LJ: Total time = 26.11097 seconds. Partitioning time = 0.0 seconds.

| thread_id | num_vertices | num_edges | barrier1_time | barrier2_time | time_taken |
|-----------|--------------|-----------|---------------|---------------|------------|
| 0 | 6510800 | 344968340 | 1.928278 | 3.455517 | 26.055284 |
| 1 | 10210820 | 344968840 | 1.920440 | 3.222096 | 26.017214 |
| 2 | 18080880 | 344968660 | 0.539212 | 2.720491 | 25.949236 |
| 3 | 62148920 | 344969620 | 0.020534 | 0.000622 | 25.949169 |